

# Premalloc:

## Pre-Runtime Heap Segmentation for Dynamic Memory Allocation

Jake Masters, Dr. Jonathan Geisler

Taylor University

### Introduction

Dynamic memory allocation involves storing and managing chunks of memory at the request of a running computer program. A program can request a memory chunk of any size at any point in time, which implies that management is non-deterministic. One of the most widely used memory allocator implementations, `dlmalloc` (based on Doug Lea's algorithm), satisfies requests for dynamic memory chunks by algorithmically populating data structures known as tree bins, which each store various numbers of same-sized chunks [1]. However, during a computer systems classroom discussion, we hypothesized that given a program's source code or runtime data, instruction overhead can be reduced. This research presents a new dynamic memory allocation algorithm we call `premalloc`.

### Background

We have limited the current scope of our research to dynamic memory allocation for Linux processes, which employ a section of the virtual address space known as the **heap**.

With respect to the algorithm, all chunks in both `dlmalloc` and `premalloc` are stored in a number of bins as illustrated below. Both implementations desire to populate all bins in such a way that they can satisfy various memory requests. In Figure 1, this end state is visualized on the right side of both arrows.

The difference between the two algorithms is `dlmalloc` begins with one chunk the size of the entire heap, while `premalloc` pre-populates each bin with what it believes is a *best-fit population* based on the information it is given. Each algorithm's initial state is given on the left side of Figure 1.

### Current State

Currently we are engaging in three threads of work:

- Defining and understanding relevant metrics
- Understanding memory allocation mechanisms
- Implementing our own memory allocator

We have completed preliminary best and worst-case analysis on the performance of `dlmalloc` using the Linux `perf` tool. We have studied the policies and mechanisms behind `dlmalloc`, which is composed of over **six thousand lines** of source code [2]. Finally, We are currently building `premalloc`, which is the software that will be our primary deliverable for this semester of research.

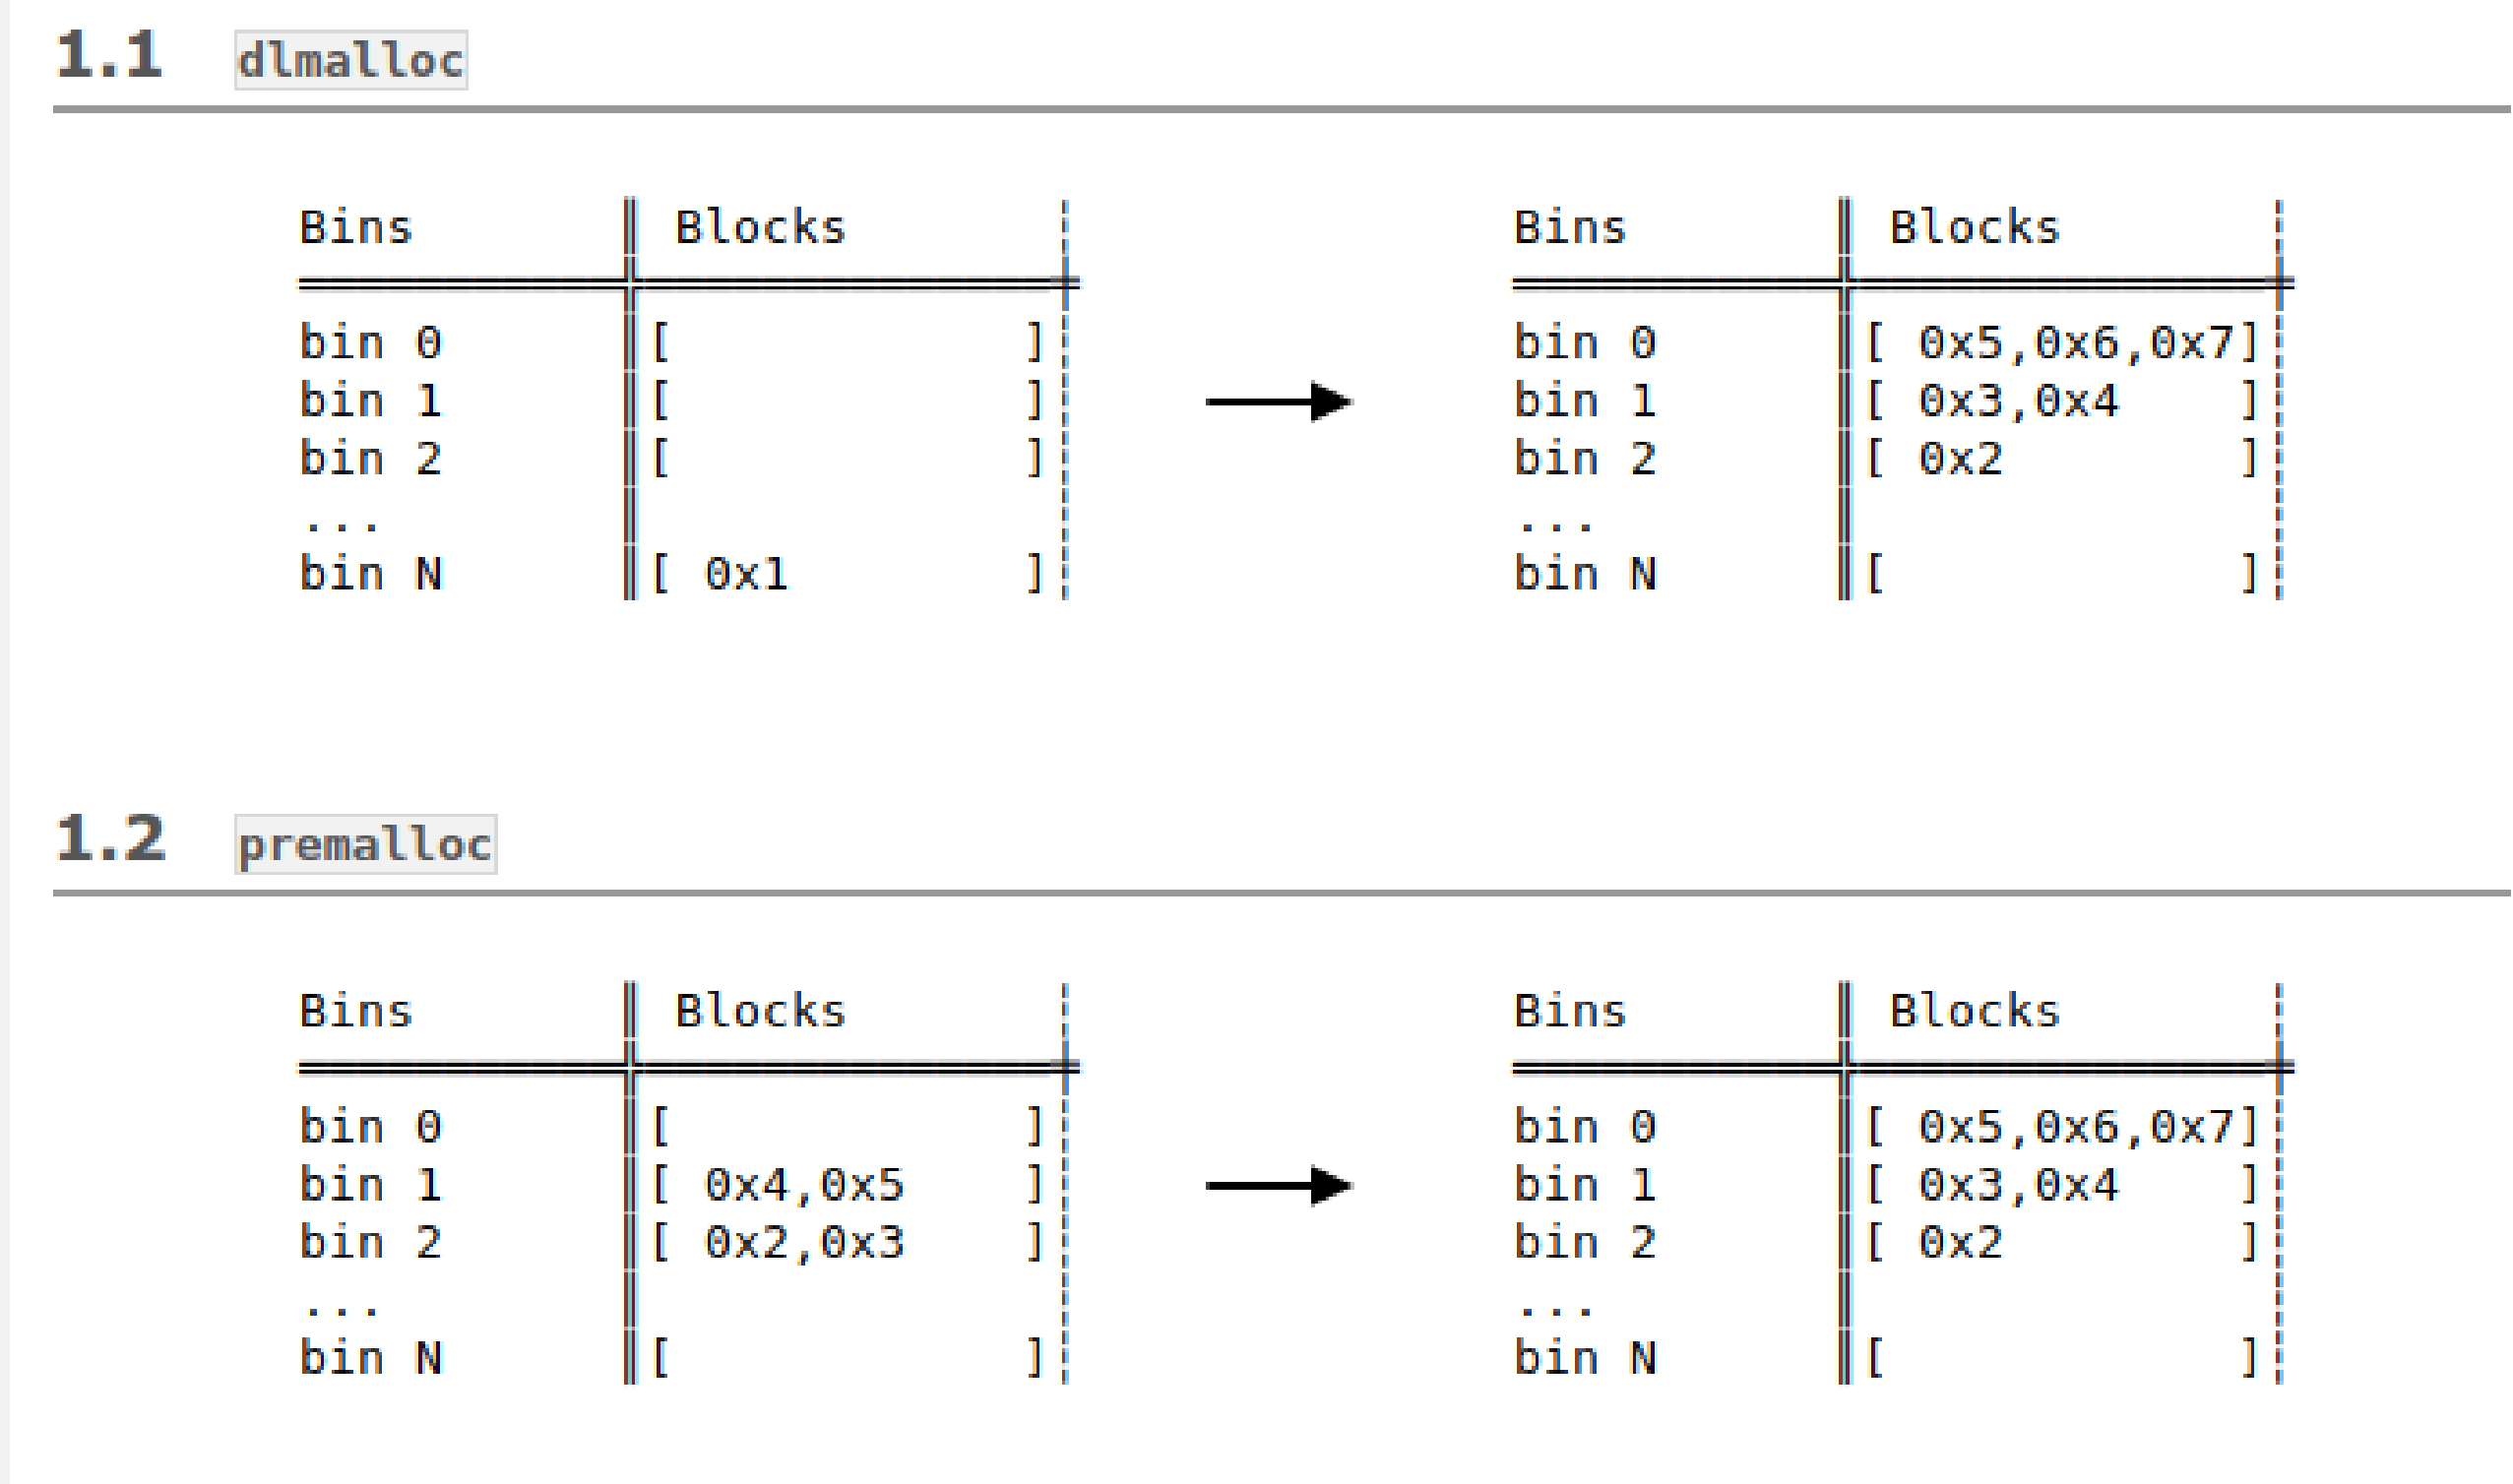


Figure 1: Both `dlmalloc` and `premalloc` need to arrive at the same end state to satisfy memory requests. The initial state for `premalloc` requires fewer instructions.

### Future Work

After our implementation of `premalloc` is in a working state, we will be conducting experiments in order to test the effectiveness of several methods in producing pre-populated free lists. Our principal metric of success is the **CPU time spent** executing memory allocation code using pre-populated free lists as opposed to algorithmically populating them at runtime. We chose this metric because it accurately measures the reduced overhead incurred by `premalloc` compared to `dlmalloc` [3]. Figure 2 illustrates the proposed process of gathering information about dynamic memory usage, in this case at runtime, to inform the memory allocator in order to pre-populate for subsequent executions.

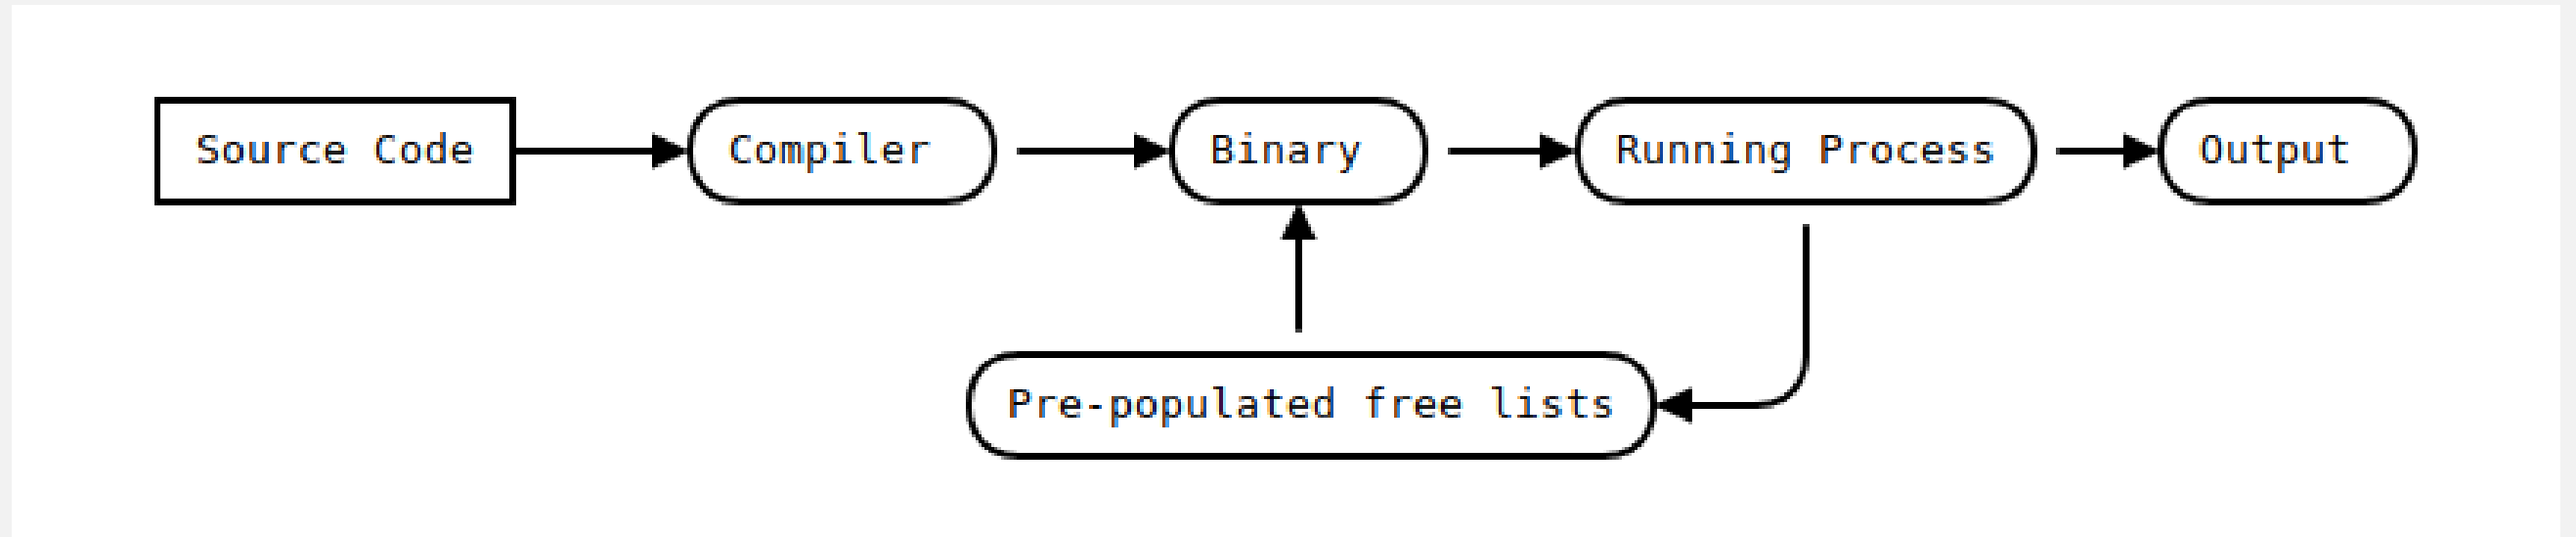


Figure 2: On the first run, runtime data is generated for a given program, which is fed back into `premalloc` on subsequent runs.

### References

- [1] P. Wilson and et al.  
Dynamic Storage Allocation: A Survey and Critical Review.
- [2] W. Fang.  
Analysis on Dynamic Memory Allocation.
- [3] E. Berger and et al.  
Reconsidering Dynamic Memory Allocation.

### Acknowledgements

We would like to thank Dr. Denning for his service and role in supervising the computer science research program in our department.